

# AIDI1009-25W: Neural Networks

## Assignment #1: Dataset Exploration and Regression

**Due Date:** 18th Feb 2025

**Student 1 Name:** Puja Shreshta (200573293)

**Student 2 Name:** Tirth Doshi (200609650)

**Student 3 Name:** Kelvin Krisna de Rodrigues Silva (2005885799)

### 1. Goals

- Understand the fundamentals of Regression Analysis.
- Explore and clean real-world datasets.
- Apply Scikit-Learn for regression modeling.
- Evaluate models using MSE & RMSE.
- Perform feature selection and hyperparameter tuning.

### 2. Introduction

#### Problem Statement

This assignment focuses on building a regression model to predict an engineering graduate's **annual salary** based on academic performance, test scores, and other factors.

#### Assumptions & Constraints

- Using `pandas`, `numpy`, `seaborn`, and `scikit-learn` for implementation.
- Missing values are handled via **imputation or column removal**.
- Data is scaled using **MinMaxScaler**.
- Performance evaluation metrics: **MSE & RMSE**.

### 3. Dataset Overview

#### About the Dataset

The dataset contains **33 independent variables** and **1 target variable (Salary)**. It provides information about students' academic backgrounds, domain knowledge, and soft skills.

#### Key Features:

- **Salary:** Target variable (Annual salary in INR).
- **10percentage, 12percentage, CollegeGPA:** Academic performance.
- **Domain Knowledge Scores:** Scores in AMCAT modules (Quant, Logical, etc.).
- **Personality Traits:** Conscientiousness, Agreeableness, etc.

## Data Challenges:

- Some fields have **missing values**.
- Not all features **strongly correlate** with Salary.
- Some categories are **imbalanced**.

## Python Code For The Problem

### 2. Perform Data Exploration

- (a) Display the **first few records** in the dataset.
- (b) Display the **number of rows and columns** of the dataset.
- (c) Display the **dataset statistics** (min, max, mean, etc.).
- (d) Display the **Null values** of each feature.
- (e) Plot **graphs** of the data to assist in data exploration.

```
In [13]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler, OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import Lasso

file_path = "Engineering_graduate_salary.csv"
df = pd.read_csv(file_path)

print("First Few Records:\n", df.head())

print("\nDataset Shape:", df.shape)

print("\nDataset Statistics:\n", df.describe())

print("\nNull Values Per Column:\n", df.isnull().sum())

plt.figure(figsize=(18, 10))
df.hist(column=['10percentage', '12percentage', 'Salary'], bins=20, figsize=(12, 10))
plt.show()
```

First Few Records:

	ID	Gender	DOB	10percentage	10board \
0	604399	f	1990-10-22	87.80	cbse
1	988334	m	1990-05-15	57.00	cbse
2	301647	m	1989-08-21	77.33	maharashtra state board,pune
3	582313	m	1991-05-04	84.30	cbse
4	339001	f	1990-10-30	82.00	cbse

	12graduation	12percentage	12board	CollegeID \
0	2009	84.00	cbse	6920
1	2010	64.50	cbse	6624
2	2007	85.17	amravati divisional board	9084
3	2009	86.00	cbse	8195
4	2008	75.00	cbse	4889

	CollegeTier	...	MechanicalEngg	ElectricalEngg	TelecomEngg	CivilEngg \
0	1	...	-1	-1	-1	-1
1	2	...	-1	-1	-1	-1
2	2	...	-1	-1	260	-1
3	1	...	-1	-1	-1	-1
4	2	...	-1	-1	-1	-1

	conscientiousness	agreeableness	extraversion	nueroticism \
0	-0.1590	0.3789	1.2396	0.14590
1	1.1336	0.0459	1.2396	0.52620
2	0.5100	-0.1232	1.5428	-0.29020
3	-0.4463	0.2124	0.3174	0.27270
4	-1.4992	-0.7473	-1.0697	0.06223

	openess_to_experience	Salary
0	0.2889	445000
1	-0.2859	110000
2	-0.2875	255000
3	0.4805	420000
4	0.1864	200000

[5 rows x 34 columns]

Dataset Shape: (2998, 34)

Dataset Statistics:

	ID	10percentage	12graduation	12percentage	CollegeID \
count	2.998000e+03	2998.000000	2998.000000	2998.000000	2998.000000
mean	6.648926e+05	77.666264	2008.080720	74.341061	5210.210807
std	3.648951e+05	10.002785	1.631814	11.120299	4776.609877
min	1.124400e+04	43.000000	1998.000000	40.000000	2.000000
25%	3.334648e+05	71.140000	2007.000000	66.000000	526.250000
50%	6.396945e+05	78.965000	2008.000000	74.000000	4027.500000
75%	9.951770e+05	85.600000	2009.000000	82.600000	8822.250000
max	1.297877e+06	97.760000	2012.000000	98.700000	18409.000000

	CollegeTier	collegeGPA	CollegeCityID	CollegeCityTier \
count	2998.000000	2998.000000	2998.000000	2998.000000
mean	1.924616	71.509857	5210.210807	0.296197

std	0.264053	8.122462	4776.609877	0.456655
min	1.000000	6.630000	2.000000	0.000000
25%	2.000000	66.530000	526.250000	0.000000
50%	2.000000	71.800000	4027.500000	0.000000
75%	2.000000	76.300000	8822.250000	1.000000
max	2.000000	99.930000	18409.000000	1.000000

	GraduationYear	...	MechanicalEngg	ElectricalEngg	TelecomEngg	\
count	2998.000000	...	2998.000000	2998.000000	2998.000000	
mean	2011.939960	...	24.138759	16.267845	31.068379	
std	36.780582	...	99.785138	86.054739	103.552963	
min	0.000000	...	-1.000000	-1.000000	-1.000000	
25%	2012.000000	...	-1.000000	-1.000000	-1.000000	
50%	2013.000000	...	-1.000000	-1.000000	-1.000000	
75%	2014.000000	...	-1.000000	-1.000000	-1.000000	
max	2017.000000	...	623.000000	660.000000	548.000000	

	CivilEngg	conscientiousness	agreeableness	extraversion	\
count	2998.000000	2998.000000	2998.000000	2998.000000	
mean	1.946965	-0.038714	0.126217	-0.008662	
std	32.241501	1.024974	0.955831	0.962695	
min	-1.000000	-3.893300	-5.781600	-4.600900	
25%	-1.000000	-0.649100	-0.435300	-0.604800	
50%	-1.000000	0.046400	0.212400	0.091400	
75%	-1.000000	0.702700	0.812800	0.672000	
max	500.000000	1.995300	1.904800	2.161700	



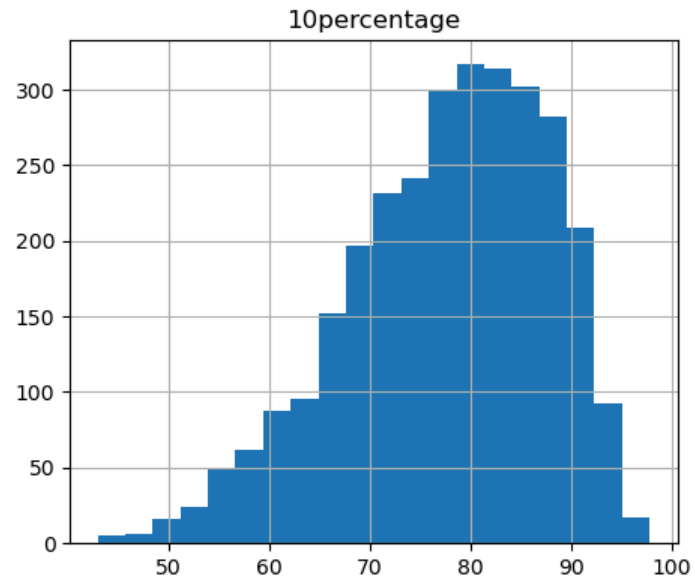
	nueroticism	openess_to_experience	Salary
count	2998.000000	2998.000000	2.998000e+03
mean	-0.145965	-0.141111	3.051748e+05
std	1.012901	1.007134	2.123312e+05
min	-2.643000	-7.375700	3.500000e+04
25%	-0.868200	-0.669200	1.800000e+05
50%	-0.172700	-0.094300	3.000000e+05
75%	0.526200	0.502400	3.700000e+05
max	3.352500	1.630200	4.000000e+06

[8 rows x 27 columns]

Null Values Per Column:

ID	0
Gender	0
DOB	0
10percentage	0
10board	0
12graduation	0
12percentage	0
12board	0
CollegeID	0
CollegeTier	0
Degree	0
Specialization	0
collegeGPA	0
CollegeCityID	0

```
CollegeCityTier      0
CollegeState         0
GraduationYear       0
English              0
Logical              0
Quant                0
Domain               0
ComputerProgramming  0
ElectronicsAndSemicon 0
ComputerScience       0
MechanicalEngg        0
ElectricalEngg        0
TelecomEngg          0
CivilEngg             0
conscientiousness     0
agreeableness         0
extraversion          0
nueroticism           0
openess_to_experience  0
Salary                0
dtype: int64
<Figure size 1800x1000 with 0 Axes>
```



### 3. Perform Initial Data Cleaning

- (a) **Delete columns** that mainly contain Null values.
- (b) **Remove duplicate columns** (obvious redundant information).
- (c) **Fill missing values** in numeric columns if necessary.
- (d) Display the **number of rows and columns** after cleaning.
- (e) Display the **features left** after cleaning.
- (f) Plot the **distribution (histogram)** of the following features:
  - **DOB**
  - **12percentage**

```
In [14]: df = df.dropna(thresh=len(df) * 0.5, axis=1)

df = df.loc[:, ~df.columns.duplicated()]

num_cols = df.select_dtypes(include=['number']).columns
imputer = SimpleImputer(strategy="mean")
df[num_cols] = imputer.fit_transform(df[num_cols])

print("\nDataset Shape After Cleaning:", df.shape)

print("\nRemaining Features:\n", df.columns)
```

Dataset Shape After Cleaning: (2998, 34)

Remaining Features:



```
Index(['ID', 'Gender', 'DOB', '10percentage', '10board', '12graduation',
      '12percentage', '12board', 'CollegeID', 'CollegeTier', 'Degree',
      'Specialization', 'collegeGPA', 'CollegeCityID', 'CollegeCityTier',
      'CollegeState', 'GraduationYear', 'English', 'Logical', 'Quant',
      'Domain', 'ComputerProgramming', 'ElectronicsAndSemicon',
      'ComputerScience', 'MechanicalEngg', 'ElectricalEngg', 'TelecomEngg',
      'CivilEngg', 'conscientiousness', 'agreeableness', 'extraversion',
      'nueroticism', 'openess_to_experience', 'Salary'],
      dtype='object')
```

```
In [15]: plt.figure(figsize=(12, 6))

plt.subplot(1, 2, 1)
df['DOB'] = pd.to_datetime(df['DOB'], errors='coerce')
df['DOB'].hist(bins=20)
plt.title("DOB Distribution")

plt.subplot(1, 2, 2)
df['12percentage'].hist(bins=20)
plt.title("12percentage Distribution")

plt.show()
```





#### 4. Analyze the Pairwise Relationship Between Features

- Use pairwise plots to visualize relationships between dataset features.

```
In [16]: selected_features = ['10percentage', '12percentage', 'Salary', 'CollegeTier']
sns.pairplot(df[selected_features])
plt.show()
```









## 5. Plot the Correlation Heatmap

- Generate a **heatmap** from the pairwise correlation matrix.

```
In [17]: numeric_df = df.select_dtypes(include=['number'])
corr_matrix = numeric_df.corr()
plt.figure(figsize=(18, 14))
sns.heatmap(numeric_df.corr(), annot=True, cmap="coolwarm", fmt=".2f", linewidths=0.6, annot_kws={"size": 8})
plt.xticks(rotation=45, ha='right', fontsize=10)
plt.yticks(fontsize=10)
plt.title("Feature Correlation Heatmap")
plt.show()
```





## 6. Perform Necessary Data Preprocessing (Transformation)

- (a) **Scale** values in numeric columns to a **(0,1) range** if needed.
- (b) **Encode categorical data** into **one-hot vectors**.
- (c) **Split the dataset** into **training, validation, and testing** sets.

```
In [18]: scaler = MinMaxScaler()

df[num_cols] = scaler.fit_transform(df[num_cols])
df['Salary'] = scaler.fit_transform(df[['Salary']])

cat_cols = df.select_dtypes(include=['object']).columns
df = pd.get_dummies(df, columns=cat_cols, drop_first=True)

corr_matrix = df.corr()

top_features = corr_matrix["Salary"].abs().sort_values(ascending=False).index[1:10]
X = df[top_features]
y = df['Salary']

X_train, X_temp, y_train, y_temp = train_test_split(X, y, test_size=0.3, random_state=42)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42)

print("\nDataset Shapes After Fixing:")
print("Training Set:", X_train.shape, y_train.shape)
print("Validation Set:", X_val.shape, y_val.shape)
print("Testing Set:", X_test.shape, y_test.shape)
```

Dataset Shapes After Fixing:  
 Training Set: (2098, 9) (2098,)  
 Validation Set: (450, 9) (450,)  
 Testing Set: (450, 9) (450,)

## 7. Train a Simple Linear Regression Model

- Use **Scikit-Learn's** `LinearRegression` model to perform regression.

## 8. Evaluate Model Performance

- Use only the features after cleaning to compute:
  - (a) **Mean Squared Error (MSE)**

### ▪ (b) Root Mean Squared Error (RMSE)

```
In [19]: model = LinearRegression()
model.fit(X_train, y_train)

y_val_pred = model.predict(X_val)
y_test_pred = model.predict(X_test)

y_val_pred = scaler.inverse_transform(y_val_pred.reshape(-1, 1))
y_test_pred = scaler.inverse_transform(y_test_pred.reshape(-1, 1))
y_val = scaler.inverse_transform(y_val.values.reshape(-1, 1))
y_test = scaler.inverse_transform(y_test.values.reshape(-1, 1))

mse_val = mean_squared_error(y_val, y_val_pred)
rmse_val = np.sqrt(mse_val)

mse_test = mean_squared_error(y_test, y_test_pred)
rmse_test = np.sqrt(mse_test)

print("\nModel Performance Metrics After Fixing:")
print("Validation Set - MSE:", mse_val, "RMSE:", rmse_val)
print("Test Set - MSE:", mse_test, "RMSE:", rmse_test)
```

Model Performance Metrics After Fixing:

Validation Set - MSE: 0.001036280594967549 RMSE: 0.0321913124144939

Test Set - MSE: 0.0036753269500475853 RMSE: 0.06062447484347873

## 10. Perform Hyperparameter Tuning

- Tune model parameters using **GridSearchCV** and compare different models.

```
In [20]: param_grid_lasso = {'alpha': [0.01, 0.1, 1, 10, 50, 100]}

lasso_model = Lasso()
grid_search_lasso = GridSearchCV(lasso_model, param_grid_lasso, cv=5, scoring='neg_mean_squared_error')
grid_search_lasso.fit(X_train, y_train)

best_lasso = grid_search_lasso.best_estimator_

y_val_pred_lasso = best_lasso.predict(X_val)
y_test_pred_lasso = best_lasso.predict(X_test)

y_val_pred_lasso = scaler.inverse_transform(y_val_pred_lasso.reshape(-1, 1))
y_test_pred_lasso = scaler.inverse_transform(y_test_pred_lasso.reshape(-1, 1))
y_val = scaler.inverse_transform(y_val.reshape(-1, 1))
y_test = scaler.inverse_transform(y_test.reshape(-1, 1))

mse_val_lasso = mean_squared_error(y_val, y_val_pred_lasso)
rmse_val_lasso = np.sqrt(mse_val_lasso)

mse_test_lasso = mean_squared_error(y_test, y_test_pred_lasso)
rmse_test_lasso = np.sqrt(mse_test_lasso)
```

```
print("\nBest Hyperparameter for Lasso Regression:", grid_search_lasso.best_params_)
print("Validation Set - MSE:", mse_val_lasso, "RMSE:", rmse_val_lasso)
print("Test Set - MSE:", mse_test_lasso, "RMSE:", rmse_test_lasso)
```

Best Hyperparameter for Lasso Regression: {'alpha': 0.01}  
 Validation Set - MSE: 0.0015594708097316746 RMSE: 0.03949013560032017  
 Test Set - MSE: 0.004226605591259178 RMSE: 0.0650123495288332

## 9. Perform Feature Selection and Repeat Step 8

- **Select key features** and recompute **MSE & RMSE** for performance comparison.

```
In [21]: model_after_selection = LinearRegression()
model_after_selection.fit(X_train, y_train)

y_test_pred_after_selection = model_after_selection.predict(X_test)
y_test_pred_after_selection = scaler.inverse_transform(y_test_pred_after_selection.reshape(-1, 1))
y_test = scaler.inverse_transform(y_test.reshape(-1, 1)) # ✅ Fix: Remove `.values`

mse_test_after_selection = mean_squared_error(y_test, y_test_pred_after_selection)
rmse_test_after_selection = np.sqrt(mse_test_after_selection)

print("\nAfter Feature Selection - Linear Regression Model:")
print("Test Set - MSE:", mse_test_after_selection, "RMSE:", rmse_test_after_selection)
```

After Feature Selection - Linear Regression Model:  
 Test Set - MSE: 0.0036753269500475853 RMSE: 0.06062447484347873

```
In [22]: feature_importance = abs(model_after_selection.coef_)
plt.figure(figsize=(10, 5))
plt.barh(X_train.columns, feature_importance)
plt.xlabel("Feature Importance")
plt.ylabel("Features")
plt.title("Feature Importance after Selection")
plt.show()
```

